

Install Libraries

In [1]: !pip install matplotlib seaborn

```
Requirement already satisfied: matplotlib in c:\users\mites\anaconda3\lib\site-packages (3.7.2)
Requirement already satisfied: seaborn in c:\users\mites\anaconda3\lib\site-packages (0.12.2)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\mites\anaconda3\lib\site-packages (from matplotlib) (1.0.5)
Requirement already satisfied: cycler>=0.10 in c:\users\mites\anaconda3\lib\site-packages (from matplotlib) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\mites\anaconda3\lib\site-packages (from matplotlib) (4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\mites\anaconda3\lib\site-packages (from matplotlib) (1.4.4)
Requirement already satisfied: numpy>=1.20 in c:\users\mites\anaconda3\lib\site-packages (from matplotlib) (1.24.3)
Requirement already satisfied: packaging>=20.0 in c:\users\mites\anaconda3\lib\site-packages (from matplotlib) (23.1)
Requirement already satisfied: pillow>=6.2.0 in c:\users\mites\anaconda3\lib\site-packages (from matplotlib) (9.4.0)
Requirement already satisfied: pyparsing<3.1,>=2.3.1 in c:\users\mites\anaconda3\lib\site-packages (from matplotlib) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\mites\anaconda3\lib\site-packages (from matplotlib) (2.8.2)
Requirement already satisfied: pandas>=0.25 in c:\users\mites\anaconda3\lib\site-packages (from seaborn) (2.0.3)
Requirement already satisfied: pytz>=2020.1 in c:\users\mites\anaconda3\lib\site-packages (from pandas>=0.25->seaborn) (2023.3.post1)
Requirement already satisfied: tzdata>=2022.1 in c:\users\mites\anaconda3\lib\site-packages (from pandas>=0.25->seaborn) (2023.3)
Requirement already satisfied: six>=1.5 in c:\users\mites\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
```

Import Libraries

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go
from mpl_toolkits.mplot3d import Axes3D
from matplotlib.colors import LogNorm
```

In [2]:

```
import warnings
warnings.filterwarnings('ignore')
```


Matplotlib Plots

a. Line Plot

A Line Plot is used to visualize trends over time or continuous data. It helps in identifying patterns, relationships, and fluctuations in data.

In [4]:


```
x = [1, 2, 3, 4, 5]
y = [10, 20, 25, 30, 40]
plt.plot(x, y, marker='o', linestyle='--', color='r')
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.title("Simple Line Plot")
plt.show()
```



b. Bar Plot


A Bar Plot is a common data visualization used to compare different categories by showing their values as rectangular bars.

```
In [5]: x = ['A', 'B', 'C', 'D']
y = [10, 20, 15, 25]
plt.bar(x, y, color='purple')
plt.xlabel("Categories")
plt.ylabel("Values")
plt.title("Bar Chart")
plt.show()
```



Horizontal Bar Plot

```
In [6]: plt.barh(x, y, color='orange')
plt.xlabel("Values")
plt.ylabel("Categories")
plt.title("Horizontal Bar Chart")
plt.show()
```




c. Scatter Plot

A scatter plot is a type of data visualization that uses dots to represent the values of two different numeric variables. It's commonly used to observe and show relationships between two continuous variables.

```
In [7]: x = [5, 7, 8, 7, 2, 17, 2, 9, 4, 11]
y = [99, 86, 87, 88, 100, 86, 103, 87, 94, 78]


plt.scatter(x, y, color='blue', marker='x')
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.title("Simple Scatter Plot")
plt.show()
```



d. Histogram

A histogram is a graphical representation of the distribution of a dataset. It consists of bars that represent the frequency of values within specified ranges (bins).

```
In [8]: data = [22, 87, 5, 43, 56, 73, 55, 54, 11, 20, 51, 5, 79, 31, 27]
plt.hist(data, bins=5, color='green', alpha=0.7)
plt.xlabel("Values")
plt.ylabel("Frequency")
plt.title("Simple Histogram")
plt.show()
```




e. Pie Chart

A pie chart is a circular statistical graphic divided into slices to illustrate proportions. Each slice represents a category's percentage of the total.

```
In [9]: sizes = [30, 25, 20, 15, 10]
labels = ['Apple', 'Banana', 'Cherry', 'Date', 'Elderberry']
colors = ['red', 'yellow', 'pink', 'brown', 'purple']

plt.pie(sizes, labels=labels, autopct='%1.1f%%', colors=colors)
plt.title("Fruit Distribution")
plt.show()
```

Fruit Distribution



f. Box Plot


A Box Plot (also called a Box-and-Whisker Plot) is a statistical visualization that represents the distribution of a dataset through five key summary statistics.

- Minimum – The smallest value in the dataset (excluding outliers).
- First Quartile (Q1) – The 25th percentile (lower quartile), where 25% of the data falls below this value.
- Median (Q2) – The 50th percentile, dividing the dataset into two equal halves.
- Third Quartile (Q3) – The 75th percentile (upper quartile), where 75% of the data falls below this value.
- Maximum – The largest value in the dataset (excluding outliers).

It is particularly useful for detecting outliers, understanding spread, and comparing distributions across categories.

```
In [10]: data = [7, 8, 9, 5, 6, 9, 8, 9, 10, 6, 7, 8]
```

```
plt.boxplot(data)
plt.title("Box Plot")
plt.show()
```




g. Error Bar Plot

An Error Bar Plot is a type of plot used to represent variability or uncertainty in data. It shows error margins around data points, indicating confidence intervals, standard deviation, or measurement errors.

In [11]:

```
x = np.arange(1, 5)
y = x ** 2
errors = [0.5, 1, 1.5, 2]

plt.errorbar(x, y, yerr=errors, fmt='o', color='black', ecolor='red', capsiz
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.title("Error Bar Plot")
plt.show()
```



h. Subplots (Multiple Plots in One Figure)

Subplots allow you to display multiple plots within a single figure in Matplotlib. This is useful when comparing multiple visualizations side by side, analyzing trends across different datasets, or presenting different perspectives of the same data.


```
In [6]: fig, axes = plt.subplots(1, 2, figsize=(10, 4))

x = np.linspace(0, 10, 100)
y1 = np.sin(x)
y2 = np.cos(x)

axes[0].plot(x, y1, color='b', label='Sine')
axes[0].set_title("Sine Wave")

axes[1].plot(x, y2, color='r', linestyle='--', label='Cosine')
axes[1].set_title("Cosine Wave")

plt.show()
```




```
In [7]: # Create the plot
plt.plot(x, y1, color='b', label='Sine')
plt.plot(x, y2, color='r', linestyle='--', label='Cosine')

# Add Labels and title
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.title("Sine and Cosine Waves")

# Add a Legend
plt.legend()

# Show the plot
plt.show()
```




i. Stacked Bar Chart

A Stacked Bar Chart is a type of bar chart used to compare parts of a whole across different categories. It represents data using rectangular bars that are stacked on top of each other, showing the contribution of each segment to the total.

```
In [9]: categories = ['A', 'B', 'C', 'D']
men = [20, 34, 30, 35]
women = [25, 32, 34, 20]

plt.bar(categories, men, label='Men', color='blue')
plt.bar(categories, women, bottom=men, label='Women', color='pink')


plt.xlabel("Categories")
plt.ylabel("Values")
plt.title("Stacked Bar Chart")
plt.legend()
plt.show()
```



```
In [15]: x = np.arange(len(categories)) # the label locations
width = 0.35 # the width of the bars

plt.bar(x - width/2, men, width, label='Men', color='blue')
plt.bar(x + width/2, women, width, label='Women', color='pink')

plt.xticks(x, categories) # Ensure x-axis labels are correct
plt.xlabel("Categories")
plt.ylabel("Values")
plt.title("Overlapping Bar Chart (Adjusted x-positions)")
plt.legend()
plt.show()
```




j. Multiple Line Plots

A Multiple Line Plot is a type of line chart that displays two or more lines on the same axes to compare trends or relationships between different datasets over a common variable, typically time or categories.

```
In [14]: x = np.linspace(0, 10, 100)
y1 = np.sin(x)
y2 = np.cos(x)

plt.plot(x, y1, label="Sine", linestyle='-')
plt.plot(x, y2, label="Cosine", linestyle='--')

plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.title("Multiple Line Plots")
plt.legend()
plt.show()
```




k. Filled Area Chart

A Filled Area Chart is a variation of a line chart where the area between the line and the x-axis (or another reference line) is filled with color to emphasize magnitude or trends over time.

In [15]:

```
x = np.linspace(0, 10, 100)
y = np.sin(x)

plt.fill_between(x, y, alpha=0.3, color='green')
plt.plot(x, y, color='black')
plt.title("Filled Area Chart")
plt.show()
```



I. Radar Chart (Spider Plot)

A Radar Chart (also called a Spider Chart, Web Chart, or Polar Chart) is a graphical method for displaying multivariate data in a two-dimensional chart with three or more quantitative variables represented on axes starting from the same central point.


1. `labels = ['A', 'B', 'C', 'D', 'E']`: This defines the labels for the axes of the radar chart. Each label represents a different category or attribute.
2. `values = [4, 3, 5, 4, 2]`: This defines the values associated with each category. These values determine how far each point is from the center of the radar chart.
3. `angles = np.linspace(0, 2 * np.pi, len(labels), endpoint=False).tolist()`: This calculates the angles at which each axis should be placed.
 - `np.linspace(0, 2 * np.pi, len(labels), endpoint=False)`: Creates evenly spaced numbers between 0 and 2π (a full circle) based on the number of labels. `endpoint=False` is important so that the last angle does not coincide with the first.
 - `.tolist()`: Converts the NumPy array to a Python list.
4. `values += values[:1] and angles += angles[:1]`: These lines are crucial for closing the radar chart. They append the first value and angle to the end of the values and angles lists. This connects the last point back to the first, creating the closed polygon shape.
5. `fig, ax = plt.subplots(figsize=(6,6), subplot_kw=dict(polar=True))`: This creates the figure and the axes for the radar chart.

- `plt.subplots(figsize=(6,6), subplot_kw=dict(polar=True))`: Creates a subplot with a polar projection (polar=True). figsize sets the size of the figure.
 - `fig`: The figure object.
 - `ax`: The axes object, which we'll use to draw on.
6. `ax.fill(angles, values, color='red', alpha=0.3)`: This fills the area enclosed by the radar chart.
- `ax.fill()`: Fills a polygon.
 - `angles`: The angles of the vertices.
 - `values`: The radial distances of the vertices.
 - `color='red'`: The fill color.
 - `alpha=0.3`: The transparency of the fill.
7. `ax.plot(angles, values, color='red')`: This draws the lines connecting the data points.
This creates the outline of the radar chart.
8. `ax.set_xticks(angles[:-1])`: Sets the positions of the x-axis ticks (the radial lines).
`angles[:-1]` is used to exclude the last angle, as it is a duplicate for closing the shape.
9. `ax.set_xticklabels(labels)`: Sets the labels for the x-axis ticks.

```
In [16]: labels = ['A', 'B', 'C', 'D', 'E']
values = [4, 3, 5, 4, 2]

angles = np.linspace(0, 2 * np.pi, len(labels), endpoint=False).tolist()
values += values[:1]
angles += angles[:1]

fig, ax = plt.subplots(figsize=(6,6), subplot_kw=dict(polar=True))
ax.fill(angles, values, color='red', alpha=0.3)
ax.plot(angles, values, color='red')
ax.set_xticks(angles[:-1])
ax.set_xticklabels(labels)
plt.title("Radar Chart")
plt.show()
```



m. Error Bars in Line Plot


Error bars represent uncertainty or variation in data points. They visually indicate the possible range of values based on standard deviation, confidence intervals, or measurement errors.

```
In [16]: errors = np.random.rand(10) * 0.5
```

```
In [6]: x = np.linspace(0, 10, 10)
y = np.sin(x)
errors = np.random.rand(10) * 0.5

plt.errorbar(x, y, yerr=errors, fmt='o-', color='black', ecolor='red', caps
plt.title("Error Bar Line Plot")
plt.show()
```


Error Bar Line Plot



n. Pie Chart with Customizations

```
In [9]: sizes = [30, 25, 20, 15, 10]
labels = ['Apple', 'Banana', 'Cherry', 'Date', 'Elderberry']
colors = ['red', 'yellow', 'pink', 'brown', 'purple']

plt.pie(sizes, labels=labels, autopct='%1.2f%%', colors=colors, startangle=
plt.title("Pie Chart with Customization")
plt.show()
```



o. 3D Scatter Plot

1. `fig = plt.figure(figsize=(8, 6))`: This line creates the figure object, which is the overall canvas for the plot. `figsize=(8, 6)` sets the size of the figure to 8 inches wide and 6 inches tall.
2. `ax = fig.add_subplot(111, projection='3d')`: This adds a subplot to the figure. 111 is a compact way of saying "1 row, 1 column, first plot". `projection='3d'` is crucial; it tells Matplotlib that this subplot should be a 3D plot. The `ax` object is what you'll use to manipulate the plot itself.
3. `x = np.random.rand(50), y = np.random.rand(50), z = np.random.rand(50)`: These lines generate 50 random numbers between 0 and 1 for each of the `x`, `y`, and `z` coordinates. `np.random.rand(50)` uses NumPy to create an array of 50 random floats.
4. `ax.scatter(x, y, z, c=z, cmap='viridis')`: This is the core of the scatter plot. `ax.scatter(x, y, z)` plots the points using the generated `x`, `y`, and `z` coordinates. `c=z` sets the color of each point based on its `z`-value. `cmap='viridis'` uses the 'viridis' colormap, which is a perceptually uniform colormap that's good for visualizing data. Other colormaps like 'plasma', 'magma', 'inferno', or 'coolwarm' could also be used.
 - Positioning: The `x`, `y`, and `z` arguments together determine the 3D location of each point. For example, the first point will be at coordinates $(x[0], y[0], z[0])$, the second at $(x[1], y[1], z[1])$, and so on.


- Coloring: The `c=z` argument tells Matplotlib to use the `z`-value of each point to determine its color. Points with higher `z`-values will have different colors than points with lower `z`-values. The `cmap='viridis'` part specifies the colormap to use for this mapping (`viridis` is a color scale that goes from purple to yellow).
5. `ax.set_xlabel("X-axis"), ax.set_ylabel("Y-axis"), ax.set_zlabel("Z-axis")`: These lines set the labels for the `x`, `y`, and `z` axes, making the plot more informative.
6. `plt.title("3D Scatter Plot")`: This sets the title of the plot.

```
In [17]: fig = plt.figure(figsize=(8, 6))
ax = fig.add_subplot(111, projection='3d')

x = np.random.rand(50)
y = np.random.rand(50)
z = np.random.rand(50)

ax.scatter(x, y, z, c=z, cmap='viridis')
ax.set_xlabel("X-axis")
ax.set_ylabel("Y-axis")
ax.set_zlabel("Z-axis")
plt.title("3D Scatter Plot")
plt.show()
```

3D Scatter Plot



p. 3D Surface Plot

A 3D Surface Plot is a three-dimensional visualization that represents continuous data as a surface in a 3D space. It is primarily used to study relationships between three numerical variables (`X`, `Y`, and `Z`).

1. `x = np.linspace(-5, 5, 50)` and `y = np.linspace(-5, 5, 50)`: These lines create two arrays, `x` and `y`, using `np.linspace()`. `np.linspace(-5, 5, 50)` generates 50 evenly spaced numbers between -5 and 5 (inclusive). These arrays will define the `x` and `y` coordinates of the grid on which the surface will be plotted.
2. `X, Y = np.meshgrid(x, y)`: `np.meshgrid()` is a crucial function here. It takes the 1D arrays `x` and `y` and creates two 2D arrays, `X` and `Y`. These `X` and `Y` arrays represent all the possible combinations of `x` and `y` values. Think of it as creating a grid of points. `X` will contain the `x`-coordinate for each point on the grid, and `Y` will contain the corresponding `y`-coordinate.
3. `Z = np.sin(np.sqrt(X2 + Y2))`: This line calculates the `z`-value for each point on the grid. It's defining a function `z = sin(sqrt(x^2 + y^2))`. NumPy's ability to work with arrays element-wise makes this calculation very efficient. **X2 and Y2** square each element of the `X` and `Y` arrays, respectively. `np.sqrt()` calculates the square root, and `np.sin()` calculates the sine. The result, `Z`, is a 2D array of `z`-values, one for each point on the grid. This `Z` array defines the height of the surface at each (`X, Y`) point.
4. `fig = plt.figure(figsize=(8, 6))` and `ax = fig.add_subplot(111, projection='3d')`: These are the same as in the scatter plot example. They create the figure and the 3D subplot where the surface will be drawn.
5. `ax.plot_surface(X, Y, Z, cmap='coolwarm', edgecolor='none')`: This is the core of the surface plot. `ax.plot_surface()` takes the `X`, `Y`, and `Z` arrays and creates the 3D surface. `cmap='coolwarm'` sets the colormap to 'coolwarm', which ranges from blue to red. `edgecolor='none'` removes the black lines that are sometimes drawn around the faces of the surface, making it look smoother.
6. `ax.set_title("3D Surface Plot")` and `plt.show()`: These lines set the title of the plot and display it, respectively.

In [19]:

```


x = np.linspace(-5, 5, 50)
y = np.linspace(-5, 5, 50)
X, Y = np.meshgrid(x, y)
Z = np.sin(np.sqrt(X**2 + Y**2))

fig = plt.figure(figsize=(8, 6))
ax = fig.add_subplot(111, projection='3d')

ax.plot_surface(X, Y, Z, cmap='coolwarm', edgecolor='none')
ax.set_title("3D Surface Plot")
plt.show()

```

3D Surface Plot



q. Contour Plot


A contour plot is a graphical representation of 3D data in two dimensions using contour lines. Each contour line represents a constant value (similar to elevation maps), allowing visualization of gradients and patterns in data.

1. `x = np.linspace(-5, 5, 50)`, `y = np.linspace(-5, 5, 50)`, `X, Y = np.meshgrid(x, y)`, and `Z = np.sin(np.sqrt(X^2 + Y^2))`: These lines are exactly the same as in the 3D surface plot code. They create the x and y coordinates and calculate the z-values using the function `z = sin(sqrt(x^2 + y^2))`.
2. `plt.contourf(X, Y, Z, levels=20, cmap='coolwarm')`: This is the core of the contour plot.
 - `plt.contourf()`: This function creates a filled contour plot. The 'f' stands for filled. If you used `plt.contour()` instead, you would get only the contour lines, without the filled colors.

- X, Y, and Z: These are the same arrays calculated earlier, providing the x, y, and z values.
 - levels=20: This specifies the number of contour levels (or the number of filled regions). The plot will divide the range of z-values into 20 intervals, and each interval will be represented by a different color. You can also provide a sequence of specific z-values for the levels instead of the number of levels.
 - cmap='coolwarm': This sets the colormap, just like in the 3D plot. It determines the colors used to fill the regions between the contour lines.
3. plt.colorbar(): This adds a colorbar to the plot. The colorbar shows the correspondence between the colors used in the contour plot and the z-values. It helps you understand what range of z-values each color represents.
4. plt.title("Contour Plot") and plt.show(): These lines set the title and display the plot.

```
In [21]: x = np.linspace(-5, 5, 50)
y = np.linspace(-5, 5, 50)
X, Y = np.meshgrid(x, y)
Z = np.sin(np.sqrt(X**2 + Y**2))

plt.contourf(X, Y, Z, levels=20, cmap='coolwarm')
plt.colorbar()
plt.title("Contour Plot")
plt.show()
```



r. Hexbin Plot (Density Visualization)

1. x = np.random.randn(1000) and y = np.random.randn(1000): These lines generate 1000 random numbers from a standard normal distribution (mean 0, standard deviation 1) for both the x and y coordinates. np.random.randn() is a convenient way to do this.


2. plt.hexbin(x, y, gridsize=30, cmap='Blues', norm=LogNorm()): This is the core of the hexbin plot.

- plt.hexbin(x, y): This function creates the hexbin plot using the x and y coordinates.
- gridsize=30: This parameter controls the resolution of the hexagonal grid. A smaller gridsize means finer hexagons (more detail), and a larger gridsize means coarser hexagons (less detail). Think of it as dividing the plot area into a grid of hexagons.
- cmap='Blues': This sets the colormap to 'Blues', meaning the hexagons will be colored with shades of blue. Denser regions (more points) will be a darker blue, and sparser regions will be a lighter blue.
- norm=LogNorm(): This is important for visualizing data with a wide range of densities. LogNorm() applies a logarithmic scaling to the color mapping. This helps to prevent very dense regions from completely dominating the color scale and making it difficult to see variations in less dense areas. If you didn't use LogNorm(), the differences in color between the less dense regions might be too subtle to perceive.

3. plt.colorbar(label='Count'): This adds a colorbar to the plot. The label='Count' argument sets the label of the colorbar to "Count," indicating that the color represents the number of points within each hexagon.

```
In [24]: x = np.random.randn(1000)
y = np.random.randn(1000)

plt.hexbin(x, y, gridsize=20, cmap='Blues', norm=LogNorm())
plt.colorbar(label='Count')
plt.title("Hexbin Plot")
plt.show()
```




s. Polar Plot

```
In [23]: theta = np.linspace(0, 2*np.pi, 100)
r = np.abs(np.sin(3*theta))


fig = plt.figure(figsize=(6,6))
ax = fig.add_subplot(111, projection='polar')

ax.plot(theta, r, color='red')
ax.set_title("Polar Plot")
plt.show()
```



t. Parallel Coordinates Plot

```
In [24]: from pandas.plotting import parallel_coordinates  
  
df = sns.load_dataset("iris")  
parallel_coordinates(df, class_column="species", colormap=plt.get_cmap("Set1"))  
plt.title("Parallel Coordinates Plot")  
plt.show()
```



u. Interactive 3D Scatter Plot (Using Plotly)

Import Library

```
In [1]: import numpy as np
```

Creating Arrays

```
In [2]: # 1D Array
arr1 = np.array([1, 2, 3, 4, 5])
print(arr1)

# 2D Array
arr2 = np.array([[1, 2, 3], [4, 5, 6]])
print(arr2)
```

```
[1 2 3 4 5]
[[1 2 3]
 [4 5 6]]
```

Using Built-in Functions

```
In [3]: # Array of zeros
zeros = np.zeros((2, 3))
print(zeros)

# Array of ones
ones = np.ones((3, 3))
print(ones)

# Array with a range of numbers
range_arr = np.arange(0, 10, 2)
print(range_arr)

# Linearly spaced array
lin_space = np.linspace(0, 1, 5)
print(lin_space)
```

```
[[0. 0. 0.]
 [0. 0. 0.]]
[[1. 1. 1.]
 [1. 1. 1.]
 [1. 1. 1.]]
[0 2 4 6 8]
[0. 0.25 0.5 0.75 1.]
```

Basic Operations

```
In [4]: arr = np.array([1, 2, 3, 4, 5])

# Addition
print(arr + 5)

# Multiplication
print(arr * 2)

# Exponentiation
print(arr ** 2)
```

```
[ 6  7  8  9 10]
[ 2  4  6  8 10]
[ 1  4  9 16 25]
```

Mathematical Functions

```
In [5]: # Sine and Cosine

angles = np.array([0, np.pi/2, np.pi])
print(np.sin(angles))
print(np.cos(angles))

# sin(pi)=0, but due to floating-point precision, it shows a very small value close to zero (1.224647e-16)

[0.000000e+00 1.000000e+00 1.2246468e-16]
[ 1.000000e+00 6.123234e-17 -1.000000e+00]
```

```
In [6]: # Square root and Exponential
```

```
nums = np.array([1, 4, 9, 16])
print(np.sqrt(nums))
print(np.exp(nums))

[1. 2. 3. 4.]
[2.71828183e+00 5.45981500e+01 8.10308393e+03 8.88611052e+06]
```

Array Indexing and Slicing

```
In [7]: #Indexing
```

```
arr = np.array([10, 20, 30, 40, 50])

# Accessing elements
print(arr[0]) # First element
print(arr[-1]) # Last element
```

```
10
50
```

```
In [8]: # Slicing
```

```
print(arr[1:4]) # Elements from index 1 to 3
print(arr[:3]) # First three elements
print(arr[::2]) # Every second element

[20 30 40]
[10 20 30]
[10 30 50]
```

```
In [9]: # 2D Array Indexing
```

```
arr2 = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])

# Accessing elements
print(arr2[1, 2]) # Row 2, Column 3

# Slicing
print(arr2[0:2, 1:3]) # Sub-matrix
```

```
6
[[2 3]
 [5 6]]
```

Broadcasting

```
In [10]: #Broadcasting
```

```
arr = np.array([1, 2, 3])
print(arr + 5) # Adds 5 to each element

[6 7 8]
```

Vectorized Operations

```
In [11]: # Vectorized Operations

a = np.array([1, 2, 3])
b = np.array([4, 5, 6])

# Element-wise multiplication
print(a * b)

[ 4 10 18]
```

Conditional Statements

```
In [12]: # Conditional Statements

arr = np.array([10, 20, 30, 40, 50])

# Boolean indexing
print(arr[arr > 30]) # Elements greater than 30

[40 50]
```

Array Manipulation

```
In [13]: # Reshape

arr = np.arange(1, 7)
reshaped_arr = arr.reshape((2, 3))
print(arr)
print(reshaped_arr)

[1 2 3 4 5 6]
[[1 2 3]
 [4 5 6]]
```

```
In [14]: # Flattening

arr2 = np.array([[1, 2, 3], [4, 5, 6]])
print(arr2)
flattened = arr2.flatten()
print(flattened)

[[1 2 3]
 [4 5 6]]
[1 2 3 4 5 6]
```

```
In [15]: # Transpose

arr3 = np.array([[1, 2, 3], [4, 5, 6]])
print(arr3)
transposed = arr3.T
print(transposed)

[[1 2 3]
 [4 5 6]]
[[1 4]
 [2 5]
 [3 6]]
```

Aggregations and Statistics

```
In [16]: # Basic Statistical Functions
```

```
arr = np.array([1, 2, 3, 4, 5])

# Sum and Mean
print(np.sum(arr))
print(np.mean(arr))

# Standard Deviation and Variance
print(np.std(arr))
print(np.var(arr))
```

```
15
3.0
1.4142135623730951
2.0
```

```
In [17]: # Min and Max
```

```
print(np.min(arr))
print(np.max(arr))
```

```
1
5
```

```
In [18]: # Axis-wise Operations
```

```
arr2 = np.array([[1, 2, 3], [4, 5, 6]])

# Sum across columns
print(np.sum(arr2, axis=0))

# Sum across rows
print(np.sum(arr2, axis=1))
```

```
[5 7 9]
[ 6 15]
```

Random Numbers and Distributions

```
In [19]: # Random integers
```

```
rand_ints = np.random.randint(0, 10, size=(2, 3))
print(rand_ints)
```

```
# Random floats between 0 and 1
rand_floats = np.random.rand(2, 3)
print(rand_floats)
```

```
# Normal Distribution
norm_dist = np.random.randn(2, 3)
print(norm_dist)
```

```
[[2 4 3]
 [0 3 8]]
[[0.98279547 0.25050344 0.2776913 ]
 [0.61946616 0.88325081 0.10157034]]
[[-0.28791975 -0.66302971 -0.62931675]
 [-1.78201822  1.083215   -0.37892021]]
```

Linear Algebra with NumPy

```
In [20]: from numpy import linalg

# Matrices
A = np.array([[1, 2], [3, 4]])
B = np.array([[5, 6], [7, 8]])

# Matrix multiplication
print(np.dot(A, B))

# Determinant
print(linalg.det(A))

# Inverse
print(linalg.inv(A))

# Eigenvalues and Eigenvectors
eigenvals, eigenvecs = linalg.eig(A)
print("Eigenvalues:", eigenvals)
print("Eigenvectors:", eigenvecs)
```

```
[[19 22]
 [43 50]]
-2.000000000000004
[[-2.  1. ]
 [ 1.5 -0.5]]
Eigenvalues: [-0.37228132  5.37228132]
Eigenvectors: [[-0.82456484 -0.41597356
 [ 0.56576746 -0.90937671]]]
```

In []:

Import Library

```
In [1]: import pandas as pd
```

Import Data

```
In [2]: cars = pd.read_csv('cars_data.csv')
```

Initial Insights

```
In [3]: cars.shape #no of rows and columns
```

```
Out[3]: (428, 15)
```

```
In [4]: rows = cars.shape[0]
columns = cars.shape[1]
```

```
In [5]: print("Rows : {}\nColumn : {}".format(rows,columns))
```

```
Rows : 428
Column : 15
```

```
In [6]: cars.head()
```

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders	Horsepower
0	Acura	MDX	SUV	Asia	All	\$36,945	\$33,337	3.5	6.0	265.0
1	Acura	RSX Type S 2dr	Sedan	Asia	Front	\$23,820	\$21,761	2.0	4.0	200.0
2	Acura	TSX 4dr	Sedan	Asia	Front	\$26,990	\$24,647	2.4	4.0	200.0
3	Acura	TL 4dr	Sedan	Asia	Front	\$33,195	\$30,299	3.2	6.0	270.0
4	Acura	3.5 RL 4dr	Sedan	Asia	Front	\$43,755	\$39,014	3.5	6.0	225.0



```
In [7]: cars.tail()
```

Out[7]:	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders	Horse
		C70 LPT								
423	Volvo	convertible 2dr	Sedan	Europe	Front	\$40,565	\$38,203		2.4	5.0
424	Volvo	C70 HPT 2dr	Sedan	Europe	Front	\$42,565	\$40,083		2.3	5.0
425	Volvo	S80 T6 4dr	Sedan	Europe	Front	\$45,210	\$42,573		2.9	6.0
426	Volvo	V40	Wagon	Europe	Front	\$26,135	\$24,641		1.9	4.0
427	Volvo	XC70	Wagon	Europe	All	\$35,145	\$33,112		2.5	5.0

◀ ▶

In [8]: `cars.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 428 entries, 0 to 427
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Make              428 non-null    object  
 1   Model             428 non-null    object  
 2   Type              428 non-null    object  
 3   Origin            428 non-null    object  
 4   DriveTrain        428 non-null    object  
 5   MSRP              428 non-null    object  
 6   Invoice           428 non-null    object  
 7   EngineSize        428 non-null    float64 
 8   Cylinders         426 non-null    float64 
 9   Horsepower        426 non-null    float64 
 10  MPG_City          428 non-null    int64   
 11  MPG_Highway       428 non-null    int64   
 12  Weight            428 non-null    int64   
 13  Wheelbase         428 non-null    int64   
 14  Length             428 non-null    int64  
dtypes: float64(3), int64(5), object(7)
memory usage: 50.3+ KB
```

In [9]: `cars.columns #name of columns`

```
Out[9]: Index(['Make', 'Model', 'Type', 'Origin', 'DriveTrain', 'MSRP', 'Invoice',
   'EngineSize', 'Cylinders', 'Horsepower', 'MPG_City', 'MPG_Highway',
   'Weight', 'Wheelbase', 'Length'],
  dtype='object')
```

In [10]: `cars.dtypes #data types`

```
Out[10]: Make          object
          Model         object
          Type          object
          Origin        object
          DriveTrain    object
          MSRP          object
          Invoice        object
          EngineSize    float64
          Cylinders     float64
          Horsepower    float64
          MPG_City      int64
          MPG_Highway   int64
          Weight         int64
          Wheelbase     int64
          Length         int64
          dtype: object
```

```
In [11]: cars.nunique() #no of unique values
```

```
Out[11]: Make          38
          Model         425
          Type          6
          Origin        3
          DriveTrain    3
          MSRP          410
          Invoice        425
          EngineSize    43
          Cylinders     7
          Horsepower    110
          MPG_City      28
          MPG_Highway   33
          Weight         348
          Wheelbase     40
          Length         67
          dtype: int64
```

```
In [12]: cars.describe()
```

	EngineSize	Cylinders	Horsepower	MPG_City	MPG_Highway	Weight	Wheelbase
count	428.000000	426.000000	426.000000	428.000000	428.000000	428.000000	428.000000
mean	3.196729	5.807512	215.854460	20.060748	26.843458	3577.953271	108.154206
std	1.108595	1.558443	72.003218	5.238218	5.741201	758.983215	8.311813
min	1.300000	3.000000	73.000000	10.000000	12.000000	1850.000000	89.000000
25%	2.375000	4.000000	165.000000	17.000000	24.000000	3104.000000	103.000000
50%	3.000000	6.000000	210.000000	19.000000	26.000000	3474.500000	107.000000
75%	3.900000	6.000000	255.000000	21.250000	29.000000	3977.750000	112.000000
max	8.300000	12.000000	500.000000	60.000000	66.000000	7190.000000	144.000000

```
In [13]: cars.isna().sum()
```

```
Out[13]:
```

Make	0
Model	0
Type	0
Origin	0
DriveTrain	0
MSRP	0
Invoice	0
EngineSize	0
Cylinders	2
Horsepower	2
MPG_City	0
MPG_Highway	0
Weight	0
Wheelbase	0
Length	0

dtype: int64

Selecting & Filtering Data

```
In [14]: cars.iloc[0:3,0:5] #iloc[index_rows, index_column]
```

```
Out[14]:
```

	Make	Model	Type	Origin	DriveTrain
0	Acura	MDX	SUV	Asia	All
1	Acura	RSX Type S 2dr	Sedan	Asia	Front
2	Acura	TSX 4dr	Sedan	Asia	Front

```
In [15]: cars.loc[0:5,'Make':'MSRP'] #Loc[values_rows, name_column]
```

```
Out[15]:
```

	Make	Model	Type	Origin	DriveTrain	MSRP
0	Acura	MDX	SUV	Asia	All	\$36,945
1	Acura	RSX Type S 2dr	Sedan	Asia	Front	\$23,820
2	Acura	TSX 4dr	Sedan	Asia	Front	\$26,990
3	Acura	TL 4dr	Sedan	Asia	Front	\$33,195
4	Acura	3.5 RL 4dr	Sedan	Asia	Front	\$43,755
5	Acura	3.5 RL w/Navigation 4dr	Sedan	Asia	Front	\$46,100

```
In [16]: cars[cars['MPG_City'] > 30] # Filter rows based on conditions
```

Out[16]:

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders
149	Honda	Civic Hybrid 4dr manual (gas/electric)	Hybrid	Asia	Front	\$20,140	\$18,451	1.4	4.0
150	Honda	Insight 2dr (gas/electric)	Hybrid	Asia	Front	\$19,110	\$17,911	2.0	3.0
154	Honda	Civic DX 2dr	Sedan	Asia	Front	\$13,270	\$12,175	1.7	4.0
155	Honda	Civic HX 2dr	Sedan	Asia	Front	\$14,170	\$12,996	1.7	4.0
156	Honda	Civic LX 4dr	Sedan	Asia	Front	\$15,850	\$14,531	1.7	4.0
159	Honda	Civic EX 4dr	Sedan	Asia	Front	\$17,750	\$16,265	1.7	4.0
352	Scion	xA 4dr hatch	Sedan	Asia	Front	\$12,965	\$12,340	1.5	4.0
353	Scion	xB	Wagon	Asia	Front	\$14,165	\$13,480	1.5	4.0
373	Toyota	Prius 4dr (gas/electric)	Hybrid	Asia	Front	\$20,510	\$18,926	1.5	4.0
379	Toyota	Corolla CE 4dr	Sedan	Asia	Front	\$14,085	\$13,065	1.8	4.0
380	Toyota	Corolla S 4dr	Sedan	Asia	Front	\$15,030	\$13,650	1.8	4.0
381	Toyota	Corolla LE 4dr	Sedan	Asia	Front	\$15,295	\$13,889	1.8	4.0
382	Toyota	Echo 2dr manual	Sedan	Asia	Front	\$10,760	\$10,144	1.5	4.0
383	Toyota	Echo 2dr auto	Sedan	Asia	Front	\$11,560	\$10,896	1.5	4.0
384	Toyota	Echo 4dr	Sedan	Asia	Front	\$11,290	\$10,642	1.5	4.0
404	Volkswagen	Jetta GLS TDI 4dr	Sedan	Europe	Front	\$21,055	\$19,638	1.9	4.0

In [17]:

```
cars[(cars['MPG_City'] > 30) & (cars['MPG_City'] < 32)]
```

Out[17]:

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders	Horsepow
353	Scion	xB	Wagon	Asia	Front	\$14,165	\$13,480	1.5	4.0	108

In [18]:

```
cars[cars['Type'] == 'Sedan'] # Filter rows with string match
```

Out[18]:

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders	Hors
1	Acura	RSX Type S 2dr	Sedan	Asia	Front	\$23,820	\$21,761	2.0	4.0	
2	Acura	TSX 4dr	Sedan	Asia	Front	\$26,990	\$24,647	2.4	4.0	
3	Acura	TL 4dr	Sedan	Asia	Front	\$33,195	\$30,299	3.2	6.0	
4	Acura	3.5 RL 4dr	Sedan	Asia	Front	\$43,755	\$39,014	3.5	6.0	
5	Acura	3.5 RL w/Navigation 4dr	Sedan	Asia	Front	\$46,100	\$41,100	3.5	6.0	
...
421	Volvo	S80 2.9 4dr	Sedan	Europe	Front	\$37,730	\$35,542	2.9	6.0	
422	Volvo	S80 2.5T 4dr	Sedan	Europe	All	\$37,885	\$35,688	2.5	5.0	
423	Volvo	C70 LPT convertible 2dr	Sedan	Europe	Front	\$40,565	\$38,203	2.4	5.0	
424	Volvo	C70 HPT convertible 2dr	Sedan	Europe	Front	\$42,565	\$40,083	2.3	5.0	
425	Volvo	S80 T6 4dr	Sedan	Europe	Front	\$45,210	\$42,573	2.9	6.0	

262 rows × 15 columns

In [19]: `cars[cars['Type'].str.contains('Sedan')]`

Out[19]:

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders	Hors
1	Acura	RSX Type S 2dr	Sedan	Asia	Front	\$23,820	\$21,761	2.0	4.0	
2	Acura	TSX 4dr	Sedan	Asia	Front	\$26,990	\$24,647	2.4	4.0	
3	Acura	TL 4dr	Sedan	Asia	Front	\$33,195	\$30,299	3.2	6.0	
4	Acura	3.5 RL 4dr	Sedan	Asia	Front	\$43,755	\$39,014	3.5	6.0	
5	Acura	3.5 RL w/Navigation 4dr	Sedan	Asia	Front	\$46,100	\$41,100	3.5	6.0	
...
421	Volvo	S80 2.9 4dr	Sedan	Europe	Front	\$37,730	\$35,542	2.9	6.0	
422	Volvo	S80 2.5T 4dr	Sedan	Europe	All	\$37,885	\$35,688	2.5	5.0	
423	Volvo	C70 LPT convertible 2dr	Sedan	Europe	Front	\$40,565	\$38,203	2.4	5.0	
424	Volvo	C70 HPT convertible 2dr	Sedan	Europe	Front	\$42,565	\$40,083	2.3	5.0	
425	Volvo	S80 T6 4dr	Sedan	Europe	Front	\$45,210	\$42,573	2.9	6.0	

262 rows × 15 columns

Drop NaN values

In [20]: `cars = cars.dropna(subset=['Horsepower']).reset_index(drop=True)`

In [21]: `cars.isna().sum()`

Out[21]:

Make	0
Model	0
Type	0
Origin	0
DriveTrain	0
MSRP	0
Invoice	0
EngineSize	0
Cylinders	2
Horsepower	0
MPG_City	0
MPG_Highway	0
Weight	0
Wheelbase	0
Length	0
dtype: int64	

In [22]: `cars.tail()`

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders	Horse
421	Volvo	C70 LPT convertible 2dr	Sedan	Europe	Front	\$40,565	\$38,203	2.4	5.0	
422	Volvo	C70 HPT convertible 2dr	Sedan	Europe	Front	\$42,565	\$40,083	2.3	5.0	
423	Volvo	S80 T6 4dr	Sedan	Europe	Front	\$45,210	\$42,573	2.9	6.0	
424	Volvo	V40	Wagon	Europe	Front	\$26,135	\$24,641	1.9	4.0	
425	Volvo	XC70	Wagon	Europe	All	\$35,145	\$33,112	2.5	5.0	

Fill NaN values in 'Cylinder' Column

In [23]: `cars[cars.Cylinders.isna()]`

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders	Horsep
245	Mazda	RX-8 4dr automatic	Sports	Asia	Rear	\$25,700	\$23,794	1.3	NaN	
246	Mazda	RX-8 4dr manual	Sports	Asia	Rear	\$27,200	\$25,179	1.3	NaN	

In [24]: `cars[cars.Make=='Mazda'].describe()`

	EngineSize	Cylinders	Horsepower	MPG_City	MPG_Highway	Weight	Wheelbase	
count	11.000000	9.000000	11.000000	11.000000	11.000000	11.000000	11.000000	11
mean	2.190909	4.444444	169.727273	21.454545	27.272727	2980.909091	105.090909	177
std	0.770006	0.881917	34.931622	3.587858	4.244783	435.576504	10.290331	13
min	1.300000	4.000000	130.000000	15.000000	19.000000	2387.000000	89.000000	156
25%	1.800000	4.000000	142.500000	18.000000	25.000000	2729.000000	103.500000	173
50%	2.000000	4.000000	160.000000	23.000000	28.000000	3029.000000	105.000000	178
75%	2.300000	4.000000	198.500000	24.000000	30.000000	3072.000000	109.000000	187
max	4.000000	6.000000	238.000000	26.000000	34.000000	3812.000000	126.000000	203

In [25]: `cars[cars.Make=='Mazda'].Cylinders.unique()`

Out[25]: `array([4., 6., nan])`

In [26]: `cars.loc[cars['Make'] == 'Mazda', ['Cylinders']]`

Out[26]: **Cylinders**

238	4.0
239	4.0
240	4.0
241	4.0
242	6.0
243	4.0
244	4.0
245	NaN
246	NaN
247	4.0
248	6.0

In [27]: `cars.loc[cars['Make'] == 'Mazda', ['Cylinders']] = cars.loc[cars['Make'] == 'Mazda'`In [28]: `cars[cars.Make=='Mazda'].isna().sum()`

Out[28]:

In [29]: `cars[cars.Make=='Mazda'].Cylinders.unique()`Out[29]: `array([4., 6.])`

Raplace Values

In [30]: `cars['MSRP'] = cars['MSRP'].str.replace('$','') # Remove $ and , in MSRP columns`
`cars['MSRP'] = cars['MSRP'].str.replace(',', '')`

In [31]: `cars['Invoice'] = cars['Invoice'].str.replace('$','') # Remove $ and , in Invoice`
`cars['Invoice'] = cars['Invoice'].str.replace(',', '')`

In [32]: `cars.head()`

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders	Horsepower	!
0	Acura	MDX	SUV	Asia	All	36945	33337	3.5	6.0	265.0	
1	Acura	RSX Type S 2dr	Sedan	Asia	Front	23820	21761	2.0	4.0	200.0	
2	Acura	TSX 4dr	Sedan	Asia	Front	26990	24647	2.4	4.0	200.0	
3	Acura	TL 4dr	Sedan	Asia	Front	33195	30299	3.2	6.0	270.0	
4	Acura	3.5 RL 4dr	Sedan	Asia	Front	43755	39014	3.5	6.0	225.0	



Change Data Type

```
In [33]: cars['MSRP'] = cars['MSRP'].astype('int') # Change data type into integer
cars['Invoice'] = cars['Invoice'].astype('int')
```

```
In [34]: cars.dtypes
```

```
Out[34]: Make          object
Model         object
Type          object
Origin        object
DriveTrain    object
MSRP          int32
Invoice       int32
EngineSize    float64
Cylinders     float64
Horsepower    float64
MPG_City      int64
MPG_Highway   int64
Weight         int64
Wheelbase     int64
Length         int64
dtype: object
```

Modifying Data

```
In [35]: cars.rename(columns={'EngineSize': 'Engine_Size'}, inplace=True) # Rename columns
```

```
In [36]: cars.columns
```

```
Out[36]: Index(['Make', 'Model', 'Type', 'Origin', 'DriveTrain', 'MSRP', 'Invoice',
               'Engine_Size', 'Cylinders', 'Horsepower', 'MPG_City', 'MPG_Highway',
               'Weight', 'Wheelbase', 'Length'],
               dtype='object')
```

```
In [37]: cars.drop(columns=['Model'], inplace=True) # Remove columns
```

```
In [38]: cars.columns
```

```
Out[38]: Index(['Make', 'Type', 'Origin', 'DriveTrain', 'MSRP', 'Invoice',
               'Engine_Size', 'Cylinders', 'Horsepower', 'MPG_City', 'MPG_Highway',
               'Weight', 'Wheelbase', 'Length'],
               dtype='object')
```

In [39]: `cars.sort_values(by='MSRP', ascending=True) # Sort DataFrame`

Out[39]:

	Make	Type	Origin	DriveTrain	MSRP	Invoice	Engine_Size	Cylinders	Horsepower
204	Kia	Sedan	Asia	Front	10280	9875	1.6	4.0	104.0
166	Hyundai	Sedan	Asia	Front	10539	10107	1.6	4.0	103.0
380	Toyota	Sedan	Asia	Front	10760	10144	1.5	4.0	108.0
343	Saturn	Sedan	USA	Front	10995	10319	2.2	4.0	140.0
205	Kia	Sedan	Asia	Front	11155	10705	1.6	4.0	104.0
...
259	Mercedes-Benz	Sedan	Europe	Rear	94820	88324	5.0	8.0	302.0
268	Mercedes-Benz	Sports	Europe	Rear	121770	113388	5.5	8.0	493.0
269	Mercedes-Benz	Sports	Europe	Rear	126670	117854	5.5	12.0	493.0
260	Mercedes-Benz	Sedan	Europe	Rear	128420	119600	5.5	12.0	493.0
332	Porsche	Sports	Europe	Rear	192465	173560	3.6	6.0	477.0

426 rows × 14 columns



Aggregation & Grouping

In [40]: `cars.groupby(['Origin', 'Type'])['MSRP'].sum()`

Out[40]:

	Origin	Type	MSRP
Asia	Hybrid		59760
	SUV		739225
	Sedan		2139813
	Sports		552681
	Truck		163069
	Wagon		254581
Europe	SUV		483460
	Sedan		3271745
	Sports		1655970
	Wagon		454215
USA	SUV		864730
	Sedan		2307495
	Sports		407315
	Truck		435524
	Wagon		156420

Name: MSRP, dtype: int32

In [41]: `cars.groupby(['Type']).agg({'MSRP': 'sum', 'MPG_City': 'mean'}) # Multiple aggregation`

Out[41]:

MSRP MPG_City

Type	MSRP	MPG_City
Hybrid	59760	55.000000
SUV	2087415	16.100000
Sedan	7719053	21.092308
Sports	2615966	18.408163
Truck	598593	16.500000
Wagon	865216	21.100000

Applying Functions

In [42]: `cars['MSRP']`

```
Out[42]: 0      36945
         1      23820
         2      26990
         3      33195
         4      43755
         ...
        421     40565
        422     42565
        423     45210
        424     26135
        425     35145
Name: MSRP, Length: 426, dtype: int32
```

In [43]: `cars['MSRP'].apply(lambda x: x**2) # Apply function to a column`

```
Out[43]: 0      73890
         1      47640
         2      53980
         3      66390
         4      87510
         ...
        421     81130
        422     85130
        423     90420
        424     52270
        425     70290
Name: MSRP, Length: 426, dtype: int64
```

In [44]: `cars['Type'].map({'Sedan': 1}) # Map specific values in a column`

```
Out[44]: 0      NaN
         1      1.0
         2      1.0
         3      1.0
         4      1.0
         ...
        421     1.0
        422     1.0
        423     1.0
        424     NaN
        425     NaN
Name: Type, Length: 426, dtype: float64
```

In [45]: `cars.apply(lambda row: row['MSRP'] - row['Invoice'], axis=1)`

```
Out[45]: 0      3608
         1      2059
         2      2343
         3      2896
         4      4741
         ...
        421     2362
        422     2482
        423     2637
        424     1494
        425     2033
Length: 426, dtype: int64
```

Merging & Joining

```
In [46]: cars_row1 = pd.read_csv('cars_data_row1.csv')
```

```
In [47]: cars_row1.shape
```

```
Out[47]: (199, 15)
```

```
In [48]: cars_row2 = pd.read_csv('cars_data_row2.csv')
```

```
In [49]: cars_row2.shape
```

```
Out[49]: (229, 15)
```

```
In [50]: cars_row = pd.concat([cars_row1, cars_row2], axis=0) # Concatenate along rows (stack)
```

```
In [51]: cars_row.shape
```

```
Out[51]: (428, 15)
```

```
In [52]: cars_col1 = pd.read_csv('cars_data_col1.csv')
```

```
In [53]: cars_col1.shape
```

```
Out[53]: (428, 7)
```

```
In [54]: cars_col2 = pd.read_csv('cars_data_col2.csv')
```

```
In [55]: cars_col2.shape
```

```
Out[55]: (428, 8)
```

```
In [56]: cars_col = pd.concat([cars_col1, cars_col2], axis=1) # Concatenate along columns
```

```
In [57]: cars_col.shape
```

```
Out[57]: (428, 15)
```

```
In [58]: cars_data_merge1 = pd.read_csv('cars_data_merge1.csv')
```

```
In [59]: cars_data_merge2 = pd.read_csv('cars_data_merge2.csv')
```

```
In [60]: cars_data_merge = cars_data_merge1.merge(cars_data_merge2, on='id', how='inner') #
```

In [61]: cars_data_merge

Out[61]:

	id	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	EngineSize	Cylinders
0	1	Acura	MDX	SUV	Asia	All	\$36,945	\$33,337	3.5	6.0
1	2	Acura	RSX Type S 2dr	Sedan	Asia	Front	\$23,820	\$21,761	2.0	4.0
2	3	Acura	TSX 4dr	Sedan	Asia	Front	\$26,990	\$24,647	2.4	4.0
3	4	Acura	TL 4dr	Sedan	Asia	Front	\$33,195	\$30,299	3.2	6.0
4	5	Acura	3.5 RL 4dr	Sedan	Asia	Front	\$43,755	\$39,014	3.5	6.0
...
423	424	Volvo	C70 LPT convertible 2dr	Sedan	Europe	Front	\$40,565	\$38,203	2.4	5.0
424	425	Volvo	C70 HPT convertible 2dr	Sedan	Europe	Front	\$42,565	\$40,083	2.3	5.0
425	426	Volvo	S80 T6 4dr	Sedan	Europe	Front	\$45,210	\$42,573	2.9	6.0
426	427	Volvo	V40	Wagon	Europe	Front	\$26,135	\$24,641	1.9	4.0
427	428	Volvo	XC70	Wagon	Europe	All	\$35,145	\$33,112	2.5	5.0

428 rows × 16 columns

Pivoting & Reshaping Data

In [62]: cars.pivot(columns='Origin', values='MSRP') #Pivot table

Out[62]:

	Origin	Asia	Europe	USA
0	36945.0	NaN	NaN	
1	23820.0	NaN	NaN	
2	26990.0	NaN	NaN	
3	33195.0	NaN	NaN	
4	43755.0	NaN	NaN	
...	
421	NaN	40565.0	NaN	
422	NaN	42565.0	NaN	
423	NaN	45210.0	NaN	
424	NaN	26135.0	NaN	
425	NaN	35145.0	NaN	

426 rows × 3 columns

In [63]: cars.T # Transpose

Out[63]:

	0	1	2	3	4	5	6	7	8	9	...	V
Make	Acura	Audi	Audi	Audi	...	V						
Type	SUV	Sedan	Sedan	Sedan	Sedan	Sedan	Sports	Sedan	Sedan	Sedan	...	Se
Origin	Asia	Europe	Europe	Europe	...	Eu						
DriveTrain	All	Front	Front	Front	Front	Front	Rear	Front	Front	Front
MSRP	36945	23820	26990	33195	43755	46100	89765	25940	35940	31840	...	31
Invoice	33337	21761	24647	30299	39014	41100	79978	23508	32506	28846	...	29
Engine_Size	3.5	2.0	2.4	3.2	3.5	3.5	3.2	1.8	1.8	3.0
Cylinders	6.0	4.0	4.0	6.0	6.0	6.0	6.0	4.0	4.0	6.0
Horsepower	265.0	200.0	200.0	270.0	225.0	225.0	290.0	170.0	170.0	220.0	...	2
MPG_City	17	24	22	20	18	18	17	22	23	20
MPG_Highway	23	31	29	28	24	24	24	31	30	28
Weight	4451	2778	3230	3575	3880	3893	3153	3252	3638	3462
Wheelbase	106	101	105	108	115	115	100	104	105	104
Length	189	172	183	186	197	197	174	179	180	179

14 rows × 426 columns

Creating Dataframe

```
In [64]: cols = ['Name', 'Gender', 'DOB', 'Mobile_no', 'Email']
rows = [['Mitesh', 'Male', '24-09-1985', '9828555308', 'mitesh.parishkar@gmail.com']]
data = pd.DataFrame(data=rows, columns=cols)
```

In [65]: data

```
Out[65]:      Name  Gender        DOB    Mobile_no           Email
0   Mitesh    Male  24-09-1985  9828555308  mitesh.parishkar@gmail.com
```

In [66]: data.dtypes

```
Out[66]: Name          object
Gender        object
DOB          object
Mobile_no     object
Email         object
dtype: object
```

Working with Dates

In [67]: import warnings

In [68]: warnings.filterwarnings('ignore')

In [69]: data['DOB'] = pd.to_datetime(data['DOB']) # Convert column to datetime

```
In [70]: data.dtypes
```

```
Out[70]: Name          object
          Gender        object
          DOB           datetime64[ns]
          Mobile_no     object
          Email         object
          dtype: object
```

```
In [71]: data['DOB'].dt.year # Extract year
```

```
Out[71]: 0    1985
          Name: DOB, dtype: int32
```

```
In [72]: data['DOB'].dt.month # Extract month
```

```
Out[72]: 0    9
          Name: DOB, dtype: int32
```

```
In [73]: data['DOB'].dt.day # Extract day
```

```
Out[73]: 0    24
          Name: DOB, dtype: int32
```

Exporting Processed Data

```
In [74]: data.to_csv('data.csv', index=False)
```

```
In [ ]:
```